

Phase 1: System Design Report

Synchronous KV Store with LRU Cache and PostgreSQL Replication

Department of Computer Engineering
Project Phase 1 Submission

Author: Monil Manish Desai

November 7, 2025

CS744 DECS - Project: HTTP-based Key-Value Server

This project implements a **distributed HTTP-based Key-Value (KV) store** with **PostgreSQL logical replication** for reliability and fault tolerance. The system uses multiple PostgreSQL instances managed through a lightweight, multi-threaded HTTP server.

GitHub Repo: <https://github.com/monil2003/cs744-project.git>

Overview

The following figure illustrates the architecture of the system.

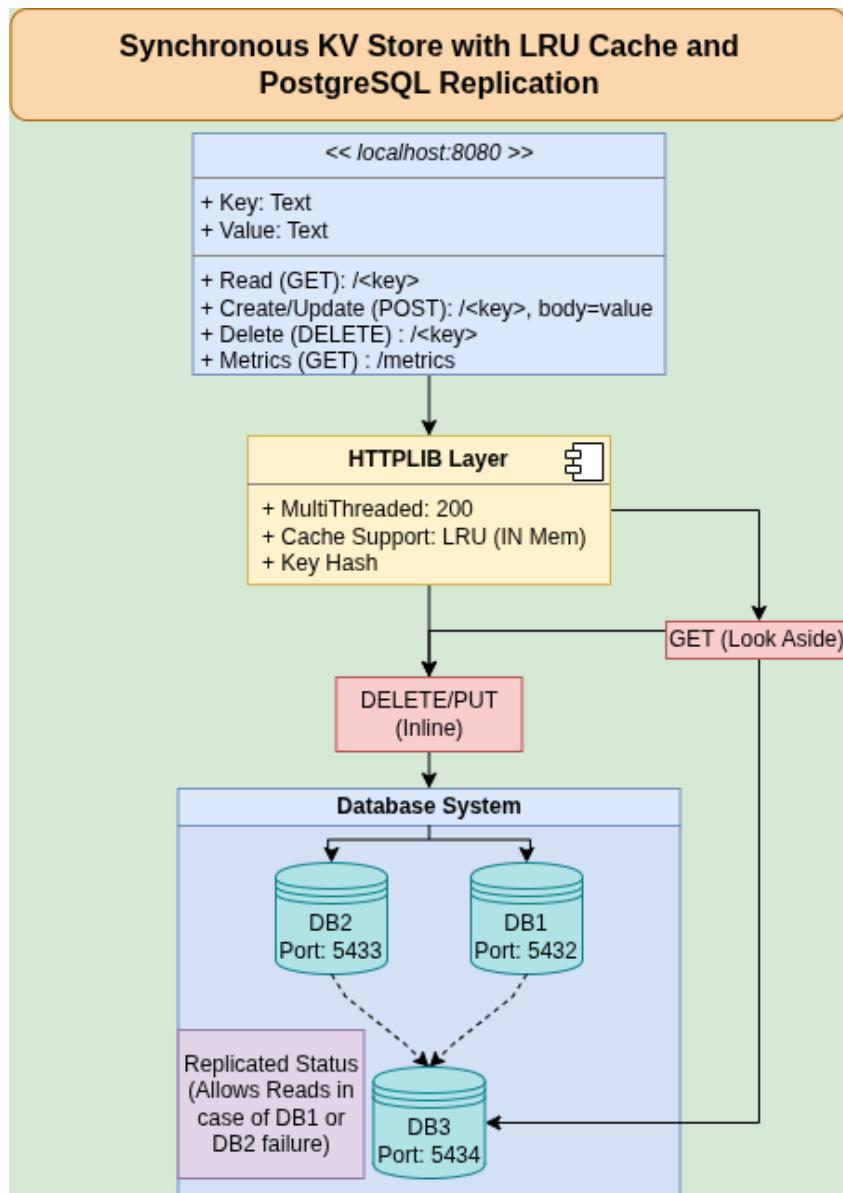


Figure 1: System Architecture of Synchronous KV Store with LRU Cache and Multi-Database Backend

HTTP Server Implementation (`httpplib`, `localhost:8080`)

The server is implemented using the `cpp-httplib` library. It listens on port **8080** and supports the following RESTful endpoints:

- `GET /<key>` – Retrieve a value for a given key.
- `POST /<key>` – Insert or update a key-value pair.
- `DELETE /<key>` – Remove a key-value pair.
- `GET /metrics` – Return cache and database performance metrics.

Each HTTP request follows one of two main execution paths:

1. **In-memory access:** When the requested key is found in the LRU cache, providing faster responses.
2. **Disk access:** When the key is not in the cache and must be retrieved or updated in the PostgreSQL database.

PostgreSQL Databases

- **db1 (port 5432)** and **db2 (port 5433)** act as primary databases.
- **db3 (port 5434)** acts as a replica subscribing to both primaries for high availability.

Replication

Logical replication ensures data consistency across nodes, with **db3** automatically synchronizing updates from both primaries.

Modes of Operation

- **Replicated Mode:** db1, db2, and db3 (high availability)
- **Direct Mode:** db1 and db2 only (for performance testing)

Features

- Multi-threaded HTTP request handling
- LRU caching for frequent lookups

- Manual PostgreSQL replication setup
- Fault-tolerant data access via db3 fallback
- Docker-based isolated deployment

Running the System

1. Start all PostgreSQL containers (db1, db2, db3).
2. Navigate to the `Databases` folder and start Docker:

```
cd Databases  
docker compose up -d
```

3. Manually configure replication as per the documentation.
4. Navigate to the server folder:

```
cd ../server
```

5. Launch the HTTP server:

```
./server_app
```

6. (Optional) Recompile using CMake:

```
cd build  
cmake --build .
```

7. Refer to the `.txt` files in each folder for example commands and code snippets.